

**Name: Suhaila Ahmed Hassan**

**Track: AI**

**Branch: Alexandria**

## Attention Is All You Need

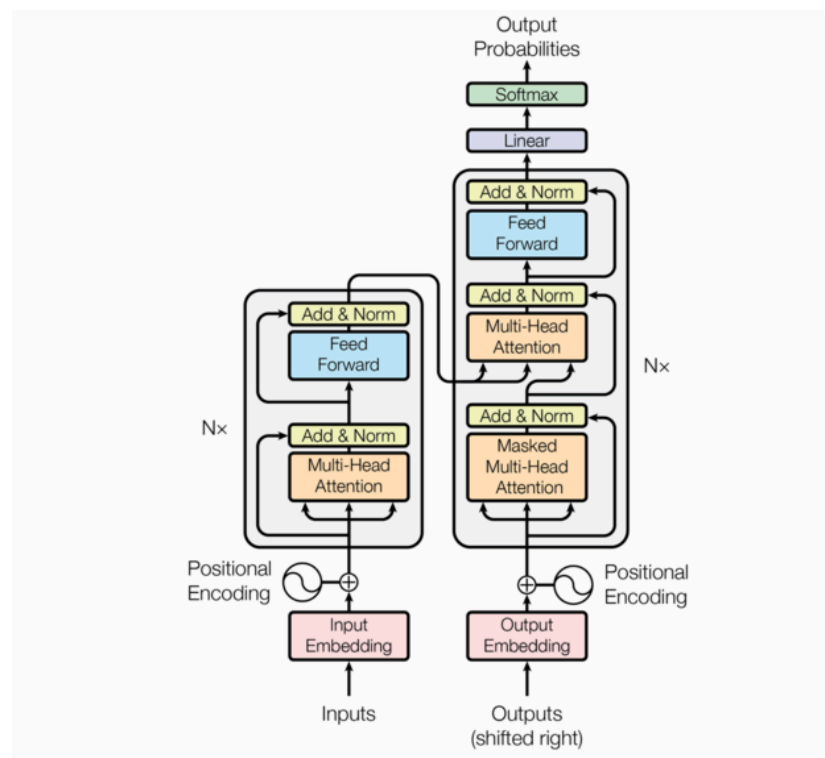
### Introduction

This paper introduces the Transformer, a novel architecture that eliminates recurrence and relies solely on attention mechanisms, enabling greater parallelization and faster training, achieving state-of-the-art results.

### Model Architecture

The Transformer is based on an encoder-decoder structure where:

- Encoder converts the input sequence into continuous representations.
- Decoder generates the output sequence one token at a time, using encoder outputs and previously generated tokens.
- Both encoder and decoder use stacked self-attention and point-wise feed-forward layers.



### 1. Encoder and Decoder Stacks

**Encoder:**

- Composed of 6 identical layers.
- Each layer has two sub-layers:
  - A. Multi-head self-attention.
  - B. Position-wise fully connected feed-forward network.
- Residual connections and layer normalization are applied to each sub-layer.


- All layers output vectors of dimension  $d_{\text{model}} = 512$ .

#### Decoder:


- Composed of 6 identical layers.
- Each layer has three sub-layers:
  - A. Masked multi-head self-attention.
  - B. Position-wise fully connected feed-forward network.
  - C. Multi-head attention over the output of the encoder stack.
- Residual connections and layer normalization are applied to each sub-layer.
- Masking prevents the decoder from seeing future positions.

## 2. Attention Mechanism

Scaled Dot-Product Attention



Multi-Head Attention



#### Attention Function:

- Maps a query and a set of key-value pairs to an output.
- Output is a weighted sum of values based on query-key similarity.

#### Scaled Dot-Product Attention:

- The input consists of queries, keys of dimension  $d_k$ , and values of dimension  $d_v$ .
- Computes attention scores using dot product of query and key, scaled by  $\sqrt{d_k}$ .
- Softmax applied to obtain attention weights.

#### Multi-Head Attention:

- Uses multiple attention "heads" in parallel ( $h = 8$ ).
- Each head has a lower dimension ( $d_k = d_v = 64$ ).
- Final output is the concatenation of all heads, followed by linear projection.
- Enables the model to attend to different positions and representation subspaces simultaneously.

#### Applications of Attention in Transformer:

The Transformer uses multi-head attention in three ways:

- Encoder-Decoder Attention: Queries come from the decoder, and keys/values from the encoder output. This lets each decoder position attend over all encoder positions, mimicking typical encoder-decoder attention in sequence-to-sequence models.

- Encoder Self-Attention: All queries, keys, and values come from the encoder's previous layer. Each encoder position attends to all positions in the previous encoder layer.
- Decoder Self-Attention: Each decoder position attends to all positions up to and including that position. To preserve the auto-regressive property, leftward information flow is blocked by masking out (setting to  $-\infty$ ) illegal connections in the softmax input of scaled dot-product attention.

### 3. Position-wise Feed-Forward Networks

- Applied independently to each position.
- Consists of two linear transformations with a ReLU activation.
- Input/output dim ( $d_{\text{model}}$ ) = 512; hidden layer dim ( $d_{\text{ff}}$ ) = 2048.

### 4. Embeddings and Softmax

- Learned embeddings convert input/output tokens to vectors of dimension  $d_{\text{model}}$ .
- In our model, same weight matrix is shared between input embeddings, output embeddings, pre-softmax linear transformation
- Embedding layers weights are multiplied by the square root of  $d_{\text{model}}$ .

### 5. Positional Encoding

- Injects information about token relative or absolute position since the model has no recurrence or convolution.
- Uses fixed sinusoidal positional encodings:
  - A.  $PE(\text{pos}, 2i) = \sin(\text{pos} / 10000^{(2i/d_{\text{model}})})$
  - B.  $PE(\text{pos}, 2i+1) = \cos(\text{pos} / 10000^{(2i/d_{\text{model}})})$
- Positional encodings are added to input embeddings.
- Fixed encodings generalize better to longer sequences than learned ones.

### Why Self-Attention

Compared to RNNs and CNNs, self-attention:

- Has lower complexity
- Allows maximum parallelism
- Has shortest path lengths for long-range dependencies
- Is more efficient when sequence length is less than hidden dimension

### Model Variation

- Attention heads: Too few or too many heads degrade performance.
- Key/Value dimension: Smaller dimensions hurt quality.
- Model size: Bigger models perform better.
- Dropout: Helps prevent overfitting.
- Positional Encoding: Learned and sinusoidal encoding performed similarly.

### Conclusion

Transformer model outperforms previous architectures on standard machine translation benchmarks.

The Transformer model, using only attention mechanisms, produces translation with better quality, has faster training and with lower cost.